



Cognitive Science 45 (2021) e13014

© 2021 The Authors. *Cognitive Science* published by Wiley Periodicals LLC on behalf of Cognitive Science Society (CSS).

ISSN: 1551-6709 online

DOI: 10.1111/cogs.13014

A Systematic Investigation of Gesture Kinematics in Evolving Manual Languages in the Lab



Wim Pouw,^{a,b}  Mark Dingemanse,^{a,c} Yasamin Motamedи,^d Aslı Özyürek^{a,b,c}

^aDonders Institute for Brain, Cognition and Behaviour, Radboud University Nijmegen

^bMax Planck Institute for Psycholinguistics, Radboud University Nijmegen

^cCenter for Language Studies, Radboud University Nijmegen

^dCentre for Language Evolution, University of Edinburgh

Received 7 July 2021; received in revised form 25 May 2021; accepted 9 June 2021

Abstract

Silent gestures consist of complex multi-articulatory movements but are now primarily studied through categorical coding of the referential gesture content. The relation of categorical linguistic content with continuous kinematics is therefore poorly understood. Here, we reanalyzed the video data from a gestural evolution experiment (Motamedи, Schouwstra, Smith, Culbertson, & Kirby, 2019), which showed increases in the systematicity of gesture content over time. We applied computer vision techniques to quantify the kinematics of the original data. Our kinematic analyses demonstrated that gestures become more efficient and less complex in their kinematics over generations of learners. We further detect the systematicity of gesture form on the level of the gesture kinematic interrelations, which directly scales with the systematicity obtained on semantic coding of the gestures. Thus, from continuous kinematics alone, we can tap into linguistic aspects that were previously only approachable through categorical coding of meaning. Finally, going beyond issues of systematicity, we show how unique gesture kinematic dialects emerged over generations as isolated chains of participants gradually diverged over iterations from other chains. We, thereby, conclude that gestures can come to embody the linguistic system at the level of interrelationships between communicative tokens, which should calibrate our theories about form and linguistic content.

Keywords: Language evolution; Silent gesture; Kinematics; Systematicity; Iterated learning

Correspondence should be sent to Wim Pouw, Montessorilaan Donders Institute for Brain, Cognition and Behaviour, Radboud University Nijmegen, 3, 6525 HR Nijmegen, The Netherlands. E-mail: w.pouw@psych.ru.nl

This is an open access article under the terms of the Creative Commons Attribution-NonCommercial License, which permits use, distribution and reproduction in any medium, provided the original work is properly cited and is not used for commercial purposes.

1. Introduction

All known natural languages combine discrete categorical elements with continuous and dynamic properties (Bolinger, 1968). For a long time, the study of human communicative behavior has focused on aspects that best yield to analysis in terms of discrete categories such as lexical items, phonological building blocks, semantic categories, and their combinatorial properties. At the same time, language use is widely acknowledged to also feature more gradient and continuous streams of behavior that do not always easily yield to an analysis in terms of discrete symbol systems (Enfield, 2009; Kendon, 2004). Here, we investigate whether kinematic measures directly derived from continuous manual movements can capture the meaning space they are designed to communicate. Using communicative silent gestures as a test case, we show how continuous movements can be studied and are patterned as evolving dynamic systems.

Manual gestures are seldomly *semiotically* studied based on the measurable part of a signal's form, namely, manual and whole-body postures in movement (i.e., the kinematic level; for exceptions see, e.g., Börstell & Lepic, 2020; Trujillo et al., 2019). Instead, gestures are mostly studied as already categorizable expressions by researchers inferring meaning from their form (McNeill, 2005). As such, the kinematics are, in one sense, reduced (from continuous to discrete tokens), and in another sense, enriched (from movement to message) with meanings that are projected onto them by human arbitrators. Here, we study more abstract aspects of gestural systems through kinematic analysis. Our aim is to show that we do not need to leave the domain of form to observe the emergence of systematic properties. We suggest that a signal's form, when studied in relation to other forms, can provide information about its linguistic properties that can complement or supplement cues that take into account conventional denotation and contextual information. Such systematicity in form—where low-level properties can serve as cues to higher-level regularities—is known from work on lexical classes in signed and spoken languages (Dingemanse, Blasi, Lupyan, Christiansen, & Monaghan, 2015; Padden et al., 2013). To observe it in existing linguistic systems, we can rely on conventional meanings and existing syntactic and semantic categories. Here, we aim to capture its emergence in rich kinematic signals as they evolve over time. In doing so, we want to contribute to an understanding of how communicative signs come to function as interrelated parts of complex dynamic systems: relatively stable ways of signaling that form higher-order structural wholes (Dale & Kello, 2018; Rączaszek-Leonardi & Kelso, 2008).

The sense-making process of individual forms becoming parts of structural wholes is essentially simulated in iterated learning experiments (Kirby, Griffiths, & Smith, 2014; Motamedi, Schouwstra, Smith, Culbertson, & Kirby, 2019). In such experiments, agents are tasked to learn a novel set of signals. These signals are iteratively transmitted to later generations and/or used in communication by later generations (iterated learning + communication), where, over many cycles of learning and use, they are affected by various transmission biases (e.g., Christiansen & Chater, 2016; Enfield, 2014). Processes of iterated learning and communication can simulate how structural properties such as systematicity, learnability, and compositionality evolve from simpler communication systems. In such simulations, communicative tokens undergo cultural evolution constrained by population dynamic properties such as historicity

(the system is constrained by past contingencies) and adaptivity (the system adapts in service of its informative goals). Such population dynamics must have played out over long temporal and vast population scales, but through these iterated learning paradigms, such processes are to some degree brought under experimental control. These evolving or emerging communicative systems can be constituted by a variety of different signal media, from simple discrete symbol sequences to more challenging continuous acoustic signals (Cornish, Dale, Kirby, & Christiansen, 2017; Ravignani, Delgado, & Kirby, 2016; Verhoef, Kirby, & de Boer, 2016).

2. Current case study

Traditionally, iterated learning studies have simulated the cultural evolution of sign systems using easily discretized word-like written forms (Scott-Phillips & Kirby, 2010). While this has made it easy to operationalize measures like compressibility, systematicity, and expressivity, it has done so at the expense of ecological validity. After all, the embodied semiotic resources that all known natural languages rely on are not carved into a predefined discrete symbol system like the Latin alphabet; instead, users and analysts of language alike must derive everything they know about linguistic systems from biological signals that are fundamentally dynamic and continuous (Pattee & Rączaszek-Leonardi, 2012).

Iterated learning work has only recently moved into the area of continuous signals. The first adaptations focused on the emergence of phonological organization in continuous acoustic signals (Verhoef, Kirby, & de Boer, 2014, 2016). Such signals do not afford a lot of semantic expressivity (but do see Ćwiek et al., 2021), so recent work has further focused on the more daunting area of continuous multi-articulator signals in the form of manual depictions or silent gestures (Motamedi et al., 2019). This experimental work studied the transmission of silent gestures created to communicate 24 concepts along two broad semantic dimensions: theme (e.g., food, religion) and function (e.g., person, location; see Fig. 1).

The two semantic dimensions provide possible axes for compression and systematization of the communicative tokens. At one extreme, one might invent 24 unique gestural utterances that are not clearly related to each other, such as in the following videos for “to sing” (<https://osf.io/d8srx/>) and “singer” (<https://osf.io/974ke/>). A more efficient encoding would be to start differentiating by functional category such that “microphone” is preceded by a general object marking gesture (<https://osf.io/r3gcp/>) and “singer” is preceded by a general person marking gesture (<https://osf.io/ex4tv/>), both followed by the same thematic marking gesture conveying “music.” Such general functional markers aid the disambiguation of related meanings, and they allow for a systematic reemployment of the same signal, thereby reducing the signaling space. In theory, once you invent four functional marker gestures and six thematic marker gestures, you can systematically recombine these to convey all 24 meanings. The communicative system then has compressed its information density from 24 to 10 information units.

Motamedi et al. (2019) indeed observed such signs of compression of the meaning space as the system developed. In early iterations of learning, fairly fine-grained iconic

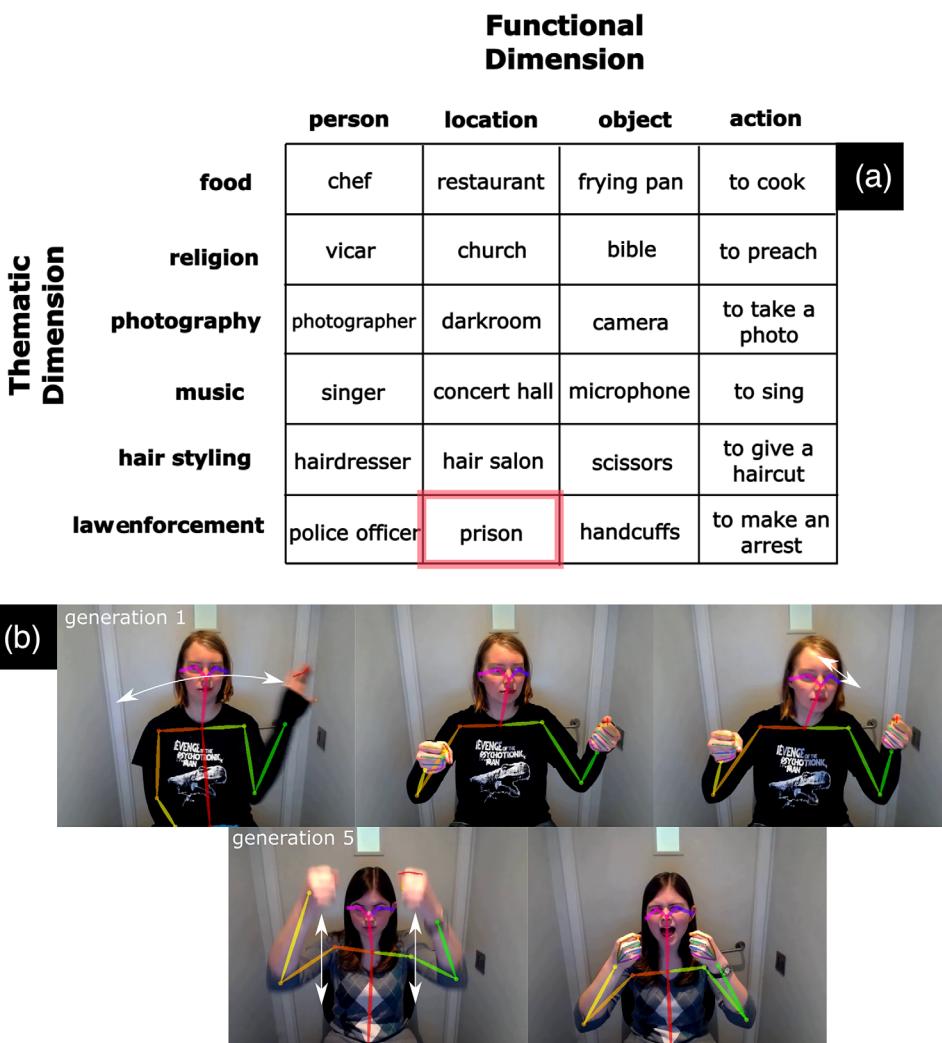


Fig. 1. Concepts to be conveyed in gesture in Motamed et al. (2019) and (motion tracking) examples (a) The concepts and categories that were used in the original experiment are shown. (b) Two example of silent gestures depicting the concept “prison.” In the first generation, a drawn out multicomponent silent gesture is produced (multiple arm movements and head movement), while in the final generation 5, a simpler gesture is produced with only two components. For the current study, we use motion tracking of the silent gestures, indicated here with a pose-skeleton overlaying the original video data.

enactments were the most common way of gesturally depicting the referents. But over generations, functional markers were found to become more prevalent and meaning components were increasingly reused across gestures. On our supplemental page, we have provided video examples (<https://osf.io/5zqnb/>) of all the gestures produced in generation 1 versus generation 5 for a particular chain (chain 1), where we highlight in red how in generation 5 there are clear

recurring functional markers for the person category (using a pointing-to-self gesture), which is absent in generation 1.

With meticulous hand coding of the different referential components of each silent gesture, Motamed et al. (2019) quantitatively tested whether there was indeed *systematicity* emerging. The gesture coding included information about the form of a particular gesture segment, such as the number of manual articulators used (1 or 2 hands), as well as the referential target of the gesture (e.g., hat; pan; turn page). Based on the full sequences of the referential components that were uniquely expressed in each gesture, *entropy* was computed, which expresses compressibility of the gesture content, i.e., the amount of information that is needed to compress the signal set.

When a lot of referential components in a gesture utterance recur between other gestures, such as in our theoretical case mentioned above, the system has a simpler structure and will show reuse of gestural components (e.g., Gibson et al., 2019). Dovetailing with the qualitative observations and other studies in this field (e.g., Verhoef et al., 2016), it was indeed found that gesture-component entropy decreased over the generations. Furthermore, the gestures were explicitly coded for the amount of marking for the functional category, and results showed that such gestures indeed occurred more often at later generations. Finally, average gesture duration—as a measure of communicative efficiency—did not reliably change over the generations, which ran counter to previous research showing a reduction in complexity over repeated gesture use (Gerwing & Bavelas, 2004; Holler & Wilkin, 2011), as well as predictions that more mature communication systems tend toward maximal efficiency (Gibson et al., 2019).

These results obtained in the lab resonate with findings from homesign (e.g., Haviland, 2013) and emerging sign languages (Senghas, Kita, & Özyürek, 2004). For example, it has been shown that in the expression of motion events, first-generation signers of Nicaraguan sign language performed more holistic presentations of path and manner, while in following generations, manner and path were segmented. Such segmentations afford novel combinatoriality and therefore increases the generativity of a language. It expresses the meaning space with fewer individual components, similar to how the participants studied by Motamed et al. (2019) started to compress the meaning space by developing ways to mark functional status across referents (e.g., “agent,” “action”).

As we see here, qualitative empirical grounding is crucial for ensuring a rich understanding of evolving multimodal signaling systems. In addition, gesture coding by human annotators makes it possible to track and quantify gestural elements encoding referents and functional dimensions, and yields measures that can be used in quantifications of entropy and the emergence of structure over time. However, it does so at the cost of reducing rich multidimensional kinematic signals to discretized sequences of coded values from a limited set. Because of this, some aspects of the evolving systems remain outside of our reach. Can we approach these developing systems in a way that is more sensitive to their continuous and dynamic properties? Do the kinematic changes that occur over generations reflect convergence on motoric “norms” or the development of unique gestural systems? In other words, do transmission chains develop their own gestural “dialects” over time? These are the questions we aim to address using kinematic methods, from a theoretical perspective influenced by dynamical systems and biosemiotics.

3. Current study

Here, we build on data from this recent iterated learning paradigm with silent gestures (Motamedi et al., 2019). With computer vision (Cao, Simon, Wei, & Sheikh, 2017) we obtained motion traces of manual and head gestures (see e.g., Lepic, Börstell, Belsitzman, & Sandler, 2016; Ripperda, Drijvers, & Holler, 2020). We then investigated kinematic inter-relationships between gestures (e.g., Beecks et al., 2015, 2016; Sato, Schouwstra, Flaherty, & Kirby, 2020), where we leverage bivariate time-series analysis (dynamic time warping (DTW)) with network topology analysis and visualization (Pouw & Dixon, 2019; Pouw et al., 2021). Through this analysis, we show that the study of gesture's form can reveal the linguistic constraints on the kinematic system as a whole.

We hypothesized that over iterations:

1. gesture kinematics become simpler;
2. gesture kinematic relationships become more systematic, and this scales with systematicity computed over gesture content coding from the original study;
3. the simplification of kinematics at the level of individual gestures is related to the systematicity of kinematic relationships across gestures;
4. idiosyncratic gesture cultures emerge as evidenced by chains drifting away from each other over time.

Importantly, the prediction that gestures simplify is based on previous research reporting simplification as judged by human annotators (e.g., Gerwing & Bavelas, 2004) and previous kinematic findings (Namboodiripad, Lenzen, Lepic, & Verhoef, 2016). We extend this research, as well as the original study (Motamedi et al., 2019), with a detailed kinematic analysis of the simplification of evolving gestures, assessing how not only salience (Namboodiripad et al., 2016) but also segmentation and temporality of gestures may change and simplify as they become part of a system of expression.

Not all kinematic changes observed will be evidence of linguistic constraints, however. Simplification over time, as in (1), could result simply from effort minimization. But when we observe increased systematicity in the system as a whole in (2) as related to kinematic simplification (3), we have a direct indication of a systematically structuring communication system. Therefore, we assess whether a Shannon-based entropy measure computed on kinematic relationships shows that the systems become more structured (i.e., compressible). Using measures that are similar to the original study, we can then also compare whether the entropy of kinematics is related to the entropy of the human-coded semantic content of the gesture. If so, we have good evidence that linguistic constraints can be objectively studied from systematic changes in gesture form. Note, again, that this is not just a methodological exercise to replicate original findings with automated methods. If we can show some equivalence between form and content analysis, then they are not on a qualitatively different analytic plane. If we can show that systematicity emerges over iterations without leaving the domain of kinematics, it means that a gesture's form is more transparent to linguistic functioning than is currently assumed.

Finally, with (4), we show how we can study chain-specific cultural evolutionary trajectories, by assessing the extent to which the communicative chains drift away from each other. The path-dependence of cultural evolution means that chains can diverge from one another over time, resulting in kinematic dialects. To assess this, we use cluster performance measures to quantify whether gestures within a chain tend to become more kinematically similar to each other and more dissimilar to gestures of other chains. This analysis is an example of the unique affordances of quantitative kinematic measures and will enable us to study, for the first time, the emergence of kinematic dialects.

4. Method

We will follow a bottom-up approach to the study of kinematics as communicative systems. In the first stage of our analysis, we demonstrate the specific changes that occur in the kinematics of the gestures. In the second stage, we assess possible systematic interrelationships in kinematic patterns of gestures through gesture network analysis (Pouw & Dixon, 2019). We will discuss each step in this procedure in the following sections and finally discuss our main gesture network entropy measure. In the supplemental information, we reserve extra space for sanity checks of automated processing and graphical descriptive results of the key measures.

4.1. Participant, design, and procedure of the original study (*Experiment 1*)

Here, we discuss the setup of the experiment, which generated the data we reanalyzed (for more detailed information, see Motamedi et al., 2019).

A seed gesture set was created with 48 pre-study participants who each depicted 1 out of 24 concepts. Thus, for each concept, there were two seed gestures performed by unique pre-study participants. Given that pre-study participants only produced one gesture, they were isolated from the other concepts that comprised the meaning space.

For the main experiment (*Experiment 1*), 50 right-handed English-speaking non-signing participants were recruited. They were allocated pairwise to one of five iteration chains. Participants were first shown a balanced subset of 24 unique seed gestures. These chain-specific seed gesture sets will be referred to as generation 0, which were followed by generations 1 through 5. In the training phase, gestures were presented in random order, and participants were asked to identify the meaning of the gesture from the 24-item meaning space, followed by feedback about their performance. They were then asked to self-record their own copy of the gesture. Participants trained with a subset of 18 items (out of 24) and completed two rounds of training.

In the testing phase, participants took turns as director and matcher to gesturally communicate (without using speech) and interpret items in the meaning space, with feedback following each trial. This director-matcher routine was repeated until both participants communicated all 24 meanings. Subsequent generations were initiated with new dyads whose training set was the gestures from one randomly selected participant from the prior generation.

The recorded videos of the seed gestures and the gesture utterances participants produced in the testing phases are the data we use here. This means that we have 50 participants conveying 24 concepts = 1200 gesture videos belonging to generations 1–5, and 48 seed gesture videos with each concept conveyed by two unique seed participants. This forms the primary data that we reanalyze using kinematic methods.

4.2. Motion tracking

Motion tracking was performed on each video recording with a sampling rate of 30 Hz. To extract movement traces, we used OpenPose (Cao et al., 2017), which is a pre-trained deep neural network approach for estimating human poses from video data (for a tutorial, see Pouw & Trujillo, 2019). We selected key points that were most likely to cover the gross variability in gestural utterances: positional x (horizontal) and y (vertical) movement traces belonging to left and right index fingers and wrists, as well as the nose. For all position traces

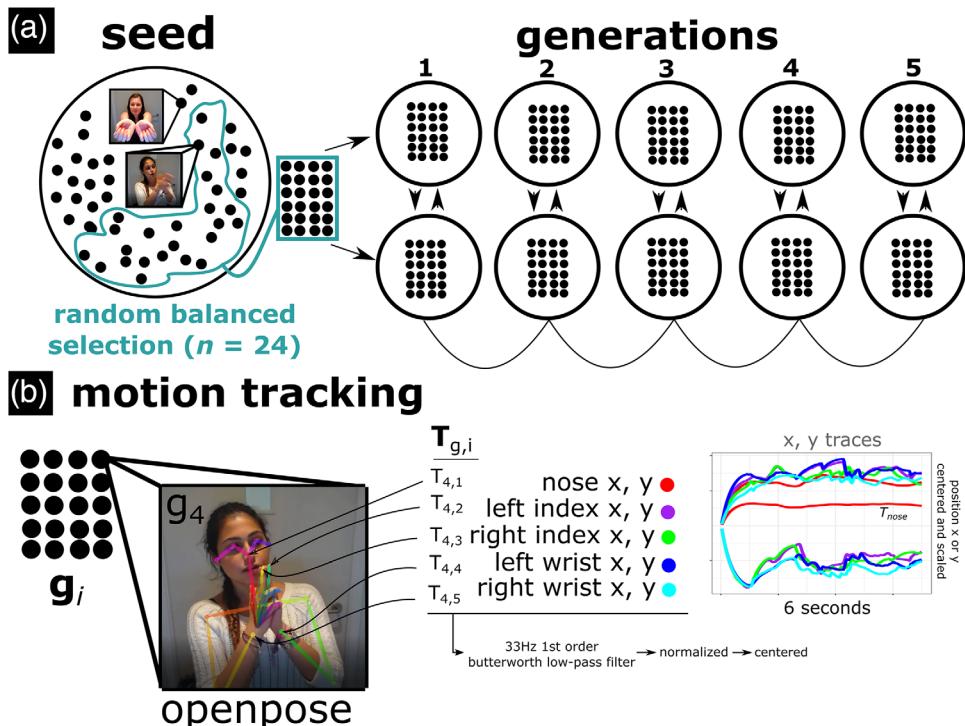


Fig. 2. Design experiment and OpenPose tracking. First steps of the general procedure (a) shows the original experiment setup (Motamedi et al., 2019), where a seed set of 24 gestures was randomly selected for each chain containing five generations. Seed gestures were used to train the first generation of each chain; after that, gestures from the previous generation were used as training data. Participants then communicated gesturally about the same concepts. (b) For our analysis, we first performed video-based motion tracking with OpenPose (Cao et al., 2017) to extract relevant two-dimensional (2D) movement traces ($T_{g,i}$) for each gesture g for body key points i (nose, the wrists, and index fingers). After motion tracking, the next steps were dynamic time warping (DTW) and gesture network analysis (Fig. 3).

and their derivatives, we applied a first-order 30 Hz low-pass Butterworth filter to smooth out high-frequency jitters having to do with sampling noise. We z-normalized and mean-centered position traces for each video to ensure that differences between subjects (e.g., body size) and within-subject differences in camera position at the start of the recording were inconsequential for our measurements. See Figure 2 for a graphical overview.

4.3. Kinematic properties

We first selected five potential measures representative of the kinematic quality of the movements in terms of segmentation, salience, and temporality, namely, submovements, intermittency, gesture space, rhythm, and temporal variability (or rhythmicity). See Fig. 3 for two example time series from which most measures can be computed. All measures were computed for each key point's time series separately and then averaged so as to get an overall score for the multimodal utterance as a whole. Based on these exploratory measures, we eventually selected three measures tracking gesture salience (*gesture space*), gesture segmentation (*intermittency score*), and gesture temporality (*temporal variability*). We discuss the motivations for selecting each measure below. Correlations between these variables and distributions are shown in supplementary materials, Fig. S1.

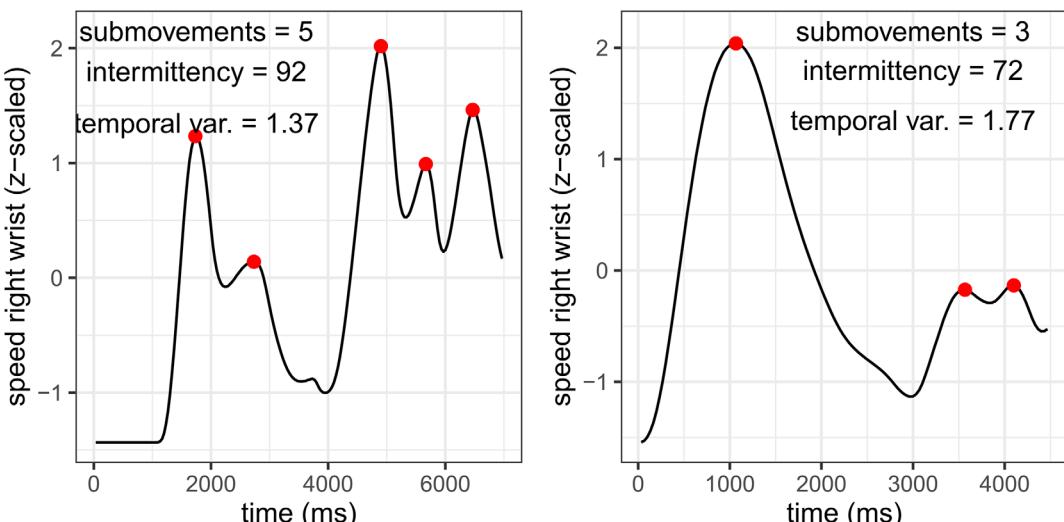


Fig. 3. Overview kinematic measures. Two time series showing right-hand wrist speed in two different trials. Our measures of segmentation and temporal variability are computed from time series like this. SEGMENTATION captures the amount of submovements (observed peaks in red), so the first time series is more segmented than the second. INTERMITTENCY captures similar information in a continuous fashion using rates of change in acceleration, yielding a higher score for the first time series than for the second. TEMPORAL VARIABILITY captures the rhythmicity of the signal and is operationalized in terms of the regularity of temporal intervals between submovements. In the first plot, red dots occur at relatively equal temporal intervals (lower temporal variability), whereas in the second, the temporal intervals are highly unequal (higher temporal variability). Finally, gesture space was calculated from the size of x,y position traces not shown here.

4.4. Gesture salience

As a measure for gesture salience or reduction, we computed a gesture space measure. This was determined by extracting the maximum vertical amplitude of a key point multiplied by the maximum horizontal amplitude, that is, the area in pixels that has been maximally covered by the movement.

4.5. Gesture segmentation

We first computed a submovement measurement similarly implemented by Trujillo, Vaitonyte, Simanova, and Özyürek (2019). Submovements are computed with a basic peak finding function which identifies and counts maxima peaks in the movement speed time series. We set the minimum inter-peak distance at eight frames, and minimum height = -1 (z-scaled; 1 std.), minimum rise = 0.1 (z-scaled).

One property of the submovement measure is that it discretizes continuous information and uses arbitrary thresholds for what counts as a submovement, thereby risking information loss about subtle intermittencies in the movement. To have a more continuous measure of intermittency (the opposite of smoothness) of the movement, we computed a dimensionless jerk measure (Hogan & Sternad, 2009). This measure is dimensionless in the sense that it is scaled by the maximum observed movement speed and duration of the movement. Dimensionless jerk is computed using the following formula:

$$\int_{t2}^{t1} x'''(t)^2 dt * \frac{D^3}{\max(v^2)}. \quad (1)$$

Here, x''' is a jerk (second derivative of the speed), which is squared and integrated over time and multiplied by duration D that is cubed over the maximum squared velocity $\max(v^2)$. We show in the supplementary materials (Fig. S2) that this measure correlates very highly with submovements; thus, we chose to only use intermittency for further analysis. Note that a *higher* intermittency score indicates more intermittent (less smooth) movement. We log-transformed our smoothness measures due to skewed distributions.

4.5.1. Gesture temporality

From the submovement measure, we computed the average interval between each submovement (in Hz), which is a measure of rhythm tempo. This measure was, as expected, highly correlated with intermittency score (see Fig. S2), as tempo goes up when more segmented movements are performed in the same time window, $r = .8$, $p < .001$, which led us to drop this measure for our analysis. Instead, we use another temporal measure that is more orthogonal to intermittency and gesture space and which captures the stability of the rhythm: the temporal variability of the movements. This measure is simply the standard deviation of the temporal interval between submovements (given in Hz): a higher score indicates more temporal variability and a lower score indicates more isochronous rhythm. Note that this measure cannot be calculated when there are less than three submovements (i.e., when there are no intervals in which we can detect the temporal variability).

4.6. Human coding and kinematic measures

For information about how these automated kinematic measures approximate hand-coded data from Motamedi et al. (2019), see Fig. S2. The human-coded data consisted of the number of unique information units of the gesture utterance, the number of repetitions in the utterance, as well as the number of segments (information units + repetitions). We should predict that our kinematic intermittency score should correlate with the number of segments, repetitions, and information units as the kinematics will have to carry those information units by contrasts in the trajectories. Fig. S2 shows the correlations for our kinematic measures and the human-coded gesture information. It shows that the number of information units (unique, repeated, or total) in the gesture as interpreted by a human coder reliably correlate with kinematic intermittency (more intermittent, more human-coded information units), gesture space (larger space, more information units), and temporal variability (more stable rhythm, more information units).

4.7. DTW

DTW is a common signal processing algorithm to quantify the similarity between temporally ordered signals (Giorgino, 2009; Mueen & Keogh, 2016; Muller, 2007). The algorithm performs a matching procedure between two time series by maximally realigning (warping) nearest values in time while preserving order and comparing their relative distances after this non-linear alignment procedure. The degree of divergence between the two time series after warping indicates how dissimilar they are. This dissimilarity is expressed with the DTW distance measure, with a higher distance score for more dissimilar time series and a lower score for more similar time series.

The time series in the current instance are multivariate, as we have a horizontal (x) and vertical (y) positional time-series data. However, DTW is easily generalizable to multivariate data and can compute its distances in a multidimensional space, yielding a multivariate dependent variant of DTW. We opt for a dependent DTW procedure here as x and y positional data are part of a single position coordinate in space. Additionally, we have six of these two-dimensional time series for each body key point. To compute a single distance measure between gestures, we computed for each gesture comparison a multivariate dependent DTW distance measure per key point, which was then summed for all key point comparisons to obtain a single distance measure D (illustrated in Fig. 5). The D measure thus reflects a general dissimilarity (higher D) or similarity (lower D) of the whole manual + head movement utterance versus another utterance. Note that the DTW procedure is applied on the entire gesture utterance, which could consist of multiple components (e.g., hand cuff gesture + pointing). This also means that gestures with a different ordering of identical components lead to high DTW distances because they are treated as two holistic gestures in the current procedure. We will come back to this in the discussion, but it should be noted that this is a drawback of the current DTW approach as it is known that flexible ordering of components is not uncommon in, for example, early developing sign languages (e.g., Ergin, Kürsat, Hartzell, Jackendoff, 2021; Ortega & Özyürek, 2020).

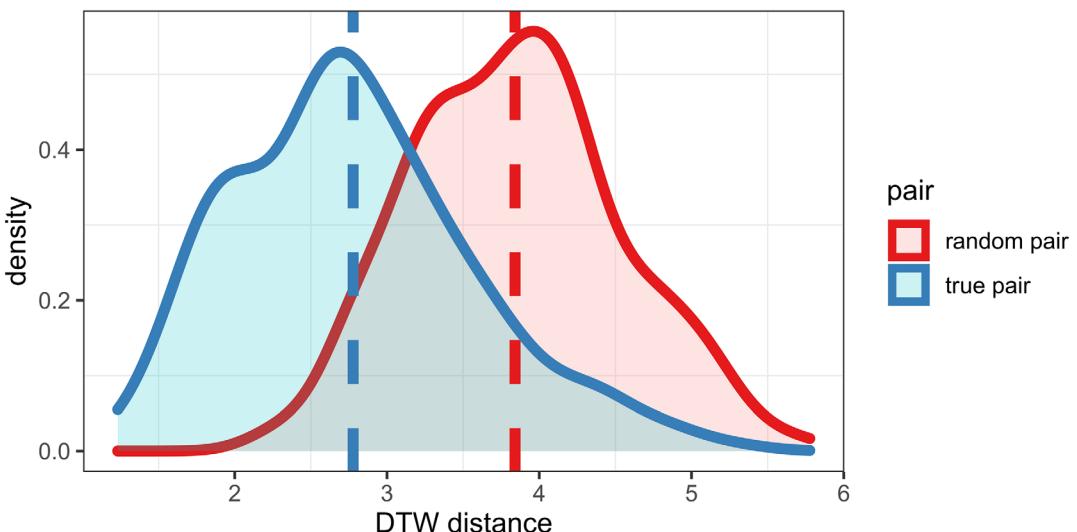


Fig. 4. Density distributions of D for true pairs and random pairs. Density distributions of D are shown for the random versus real pairs. With D based on head, wrist, and finger movements, there is good discriminability between real versus falsely paired gestures, confirming that our approach is tracking gesture similarity well.

We used the R package “DTW” (Giorgino, 2009) to produce the multivariate distances per key point. The DTW distance measure was normalized for both time series’ lengths, such that average distances are expressed per unit time rather than summing distances over time which would yield higher (and biased) distance estimates for longer time series (i.e., longer gesture videos). For a further conceptual overview and methodological considerations of our DTW procedure, see Pouw and Dixon (2019).

As a demonstration that our D measure reflects actual differences in kinematics, we computed for each individual in each chain the difference between a gesture seed and the gesture that the individual produced to copy it for generation 1. These “true pairs” must be maximally similar (lower D) as the individual produced their copied gesture shortly after first exposure in the training phase, which should lead to high faithfulness in reproduction. We contrast this with a false or random comparison of the same gesture in generation 1 with a gesture seed that was neither in the same functional nor thematic category. These false random pairs must be more dissimilar and should produce higher DTW distances. Fig. 4 shows the distributions of the distances observed. DTW distance distributions were reliably different, $t(469.77) = 15.82$, $p < .001$, Cohen’s $d = 1.44$, for the true pair, $M = 2.78$ ($SD = 0.78$), as compared to the random pair, $M = 3.84$ ($SD = 0.69$).

We find that adding head movement trajectory to our D calculation significantly increases false-real pair discriminability in comparison when we compute our D measure on only manual key points (left/right wrist and index fingers), change in Cohen’s $d = 0.41$, change D real versus false = 0.33, $p < .001$. Therefore, we conclude that in the current experiment, the gesture utterances are also crucially defined by head movements as well. This is a novel finding in and of itself and demonstrates the multi-articulatory nature of silent gestures.

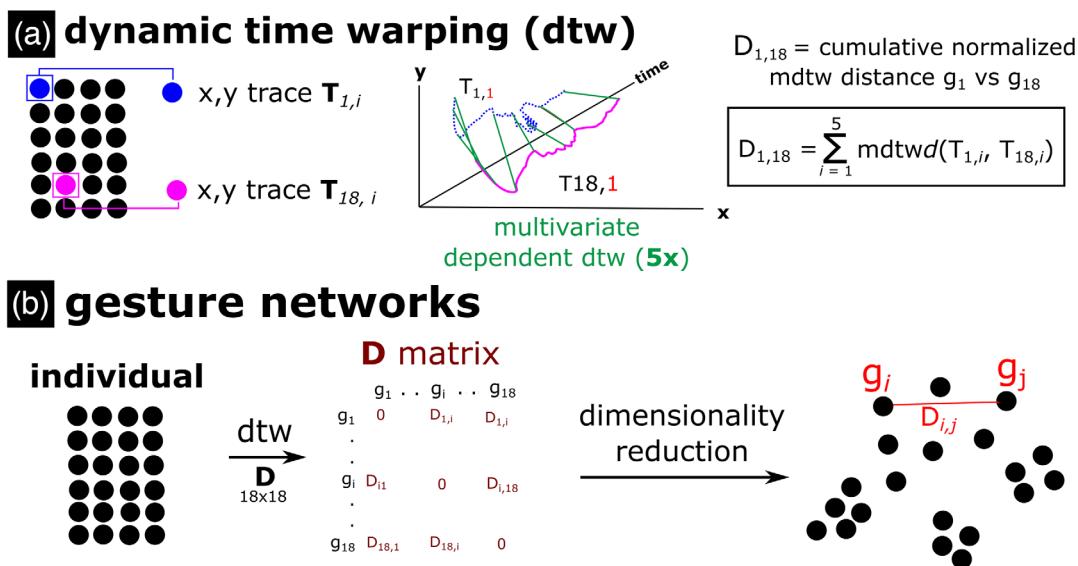


Fig. 5. General method gesture network analysis. (a) For each body part in a gesture comparison, we used DTW to compute a multivariate normalized distance gesture, which we summed into an overall distance measure D for each gesture comparison within a gesture set. (b) All distance measures were saved into a distance matrix \mathbf{D} containing all gesture comparisons $D_{i,j}$ within the comparison set, resulting in a 24×24 distance matrix. The distance matrix can be visualized as a weighted graph through dimensionality reduction techniques, such that nodes indicate gesture utterances and the distance (or weight) between gesture nodes representing the ' D ' measure, indicating dissimilarity.

4.8. Gesture kinematic networks

Graphically shown in Fig. 5, we constructed for each participant (nested in generation and chain), as well as each seed gesture set (seed set belonging to that chain), a distance matrix \mathbf{D} , containing the continuous D comparisons for each gesture $D_{i,j}$ produced by that participant with each other gesture produced by that participant, yielding a 24×24 matrix. The diagonal contains zeros for gesture comparisons that are identical ($D_{i,j} = 0 | i = j$). These characteristics make \mathbf{D} a weighted symmetric distance matrix.

For each distance matrix, we can construct a visual geometric representation of its topology by projecting the distance of gesture tokens on a 2D plane using a dimensionality reduction technique called “t-SNE,” a variant of stochastic neighbor embedding (Maaten & Hinton, 2008). These 2D representations show locations of gesture nodes, with distances between gesture nodes approximating our D measure. Note though that these 2D approximations in the case of t-SNE are exaggerated projections of the data and should be distinguished from the actual high-dimensional structure of the data. The uncompressed distance matrices are used to calculate entropy and other measures. We refer to these measurements as “network properties,” as these measures are intuitively understood in the network or geometric terms. For calculations of network entropy, we use the R package “igraph” (Csárdi, 2019), and for dimensionality reduction, we use R package “tsne” (Donaldson, 2016). On our supplemental

page, we again show video examples of all the gestures produced in generation 1 versus generation 5 for a particular chain (chain 1) but now with videos spatially located according to their coordinates in kinematic space (<https://osf.io/wbmf9/>). Examples are highlighted in red, where kinematic similarity increases from generation 1 to 5 due to functional markers being used for the category “location.”

4.8.1. Kinematic entropy

Entropy is a measure that quantifies the compressibility of data structures and has been used to gauge the combinatorial structure of communicative tokens in the field of language evolution (e.g., Verhoef et al., 2016; for theoretical grounding, see Gibson et al., 2019). In the original experiment, Motamedi et al. (2019) computed entropy from the gesture content codings, which captured recurrent information units between gestures. In our case, entropy quantifies the degree to which there are similar or more diverse edge lengths (i.e., similar/diverse levels of dissimilarity “D” between combinations of two gesture trajectories). If they are more similar, lower entropy reflects that communicative tokens relate to each other in more structural ways. So our measure of network entropy gauges how compressible kinematic interrelationships are, which is conceptually related to the systematic recurrence of information units between the human judged gesture content.

The network entropy measure we used (see Eagle, Macy, & Claxton, 2010) is almost identical to a classic Shannon entropy calculation used in the original study to quantify the systematicity of the gesture’s content (Motamedi et al., 2019), where $Entropy H(X) = - \sum p(X) \log p(X)$. The only difference is that our measure is computed on the distances for each node relative to the shortest path to the other nodes and then normalized by the number of gesture distances. So our measure quantifies the topological diversity of the gesture relationships, where a lower score indicates more similar relationships and a higher score indicates a more randomly distributed set of relationships. Specifically, for each gesture node, we compute the diversity of kinematic distances to other gestures, using a scaled Shannon entropy measure:

$$H(i) = - \sum_{j=1}^k p_{ij} \log p_{ij} / \log(k_i). \quad (2)$$

Here, k_i is the number of gesture connections for gesture i , and p_{ij} is the proportional distance:

$$p_{ij} = D_{ij} / \sum_{j=1}^k D_{ij}. \quad (3)$$

Here, p_{ij} is the distance between gesture i and gesture j divided by the total distance involving gesture i . Fig. 6 shows a graphical example of different network structures and the concomitant entropy measure.

Gesture kinematic culture. To assess whether there is a kinematic culture emerging such that gestures in a specific chain are over the generations becoming more similar in kinematics as compared to gestures from another chain, we leverage cluster performance measures. For

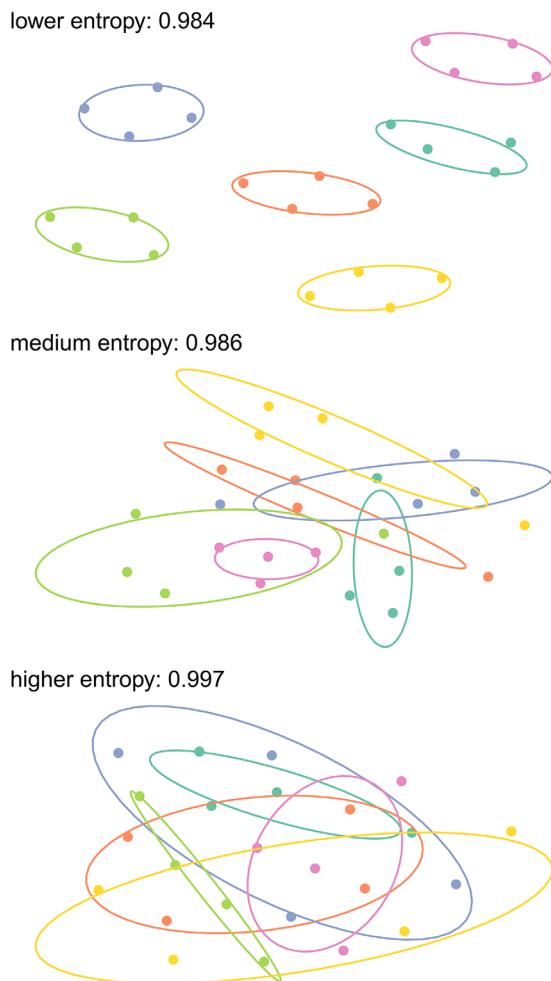


Fig. 6. Example network entropy. Simulated data showing six clusters with low variance in distances (top panel), higher variance (middle panel), or randomly distributed distances (lower panel). More variable and random distributions of node distances yield higher entropy scores. In contrast, entropy is lowest when interrelationships are distributed in a more systematic way (top panel).

each generation, we assess whether gestures from a particular chain are also likely to cluster in a super-ordinate kinematic space that includes all gestures performed across generations (i.e., gestures produced in chains 1 through 5). Clustering can be quantified in several ways. In our analysis, we report on two well-known cluster performance measures: Dunn index and Silhouette width (Yadav, Tomar, & Agarwal, 2013). In general, cluster performance measures relate within-cluster distances between nodes (minimal when clusters are stable) to between-cluster distances (maximal when clusters are stable), though they vary in how they compute the within and between distances. The Dunn index quantifies the compactness of the clusters assigned (chains in our case) and relates the minimum distance between centroids of each

cluster to maximal distance between points, where higher values indicate better clustering. However, this Dunn index measure only yields five data points in our case, one for each chain, which makes it hard to perform a statistical test. Therefore, we will also compute a token level measure of Silhouette width, which, for each token, relates the mean distance to other tokens within its cluster to the minimum distance between a member of a neighboring cluster.

5. Main results

We first report changes in kinematic features over generations. Then we consider the change in relations between communicative tokens over generations as indexed by kinematic network entropy. We also relate kinematic changes to network-level changes. Finally, we consider how chains diverge over time, allowing a peek into the emergence of unique gesture cultures.

5.1. Kinematic features

A key aim of our analysis is to capture the fine-grained kinematic features that drive changes in the gestural systems over generations, which are hard to capture with a manual coding system focusing on the semiotic relation between gesture and meaning. All three of our kinematic measures show the hallmarks of increased communicative efficiency through reduced kinematic complexity over generations (Fig. 7).

We performed mixed effects regression analysis for assessing potential kinematic changes as a function of generation, with random intercept for participants nested within chains (random slopes did not converge). Generation reliably predicted intermittency of the movements relative to a base model (chi-squared change (1) = 76.66, $p < .001$, model $R^2 = 0.06$). In this model, generation predicted lower intermittency score (b estimate = -0.2263, $t(1135.00) = -8.90$, $p < .001$, Cohen's $d = -0.53$). We also observe lower temporal variability as a function of generation (chi-squared change (1) = 24.12, $p < .001$, model $R^2 = 0.05$), indicating more stable rhythmic movements at later generations (b estimate = -0.0693, $t(332.00) = -4.97$, $p < .001$, Cohen's $d = -0.55$).

Finally, over generations gesture space decreased (chi-squared change (1) = 24.45, $p < .001$). Model estimated gesture space was less for later generations (b estimate = -2.2100, $t(1174.00) = -4.97$, $p < .001$, Cohen's $d = -0.29$).

Subtle changes in kinematic features are hard to capture using human coding, and indeed the rough proxies for this used by Motamed et al. (2019; length and number of repetitions of coded information units) did not demonstrate increased communicative efficiency. Here, we are able to capture increased efficiency by quantifying fine-grained kinematic features at the level of gesture tokens. Using independently motivated measures, we found that gestures were on average smaller, less temporally variable, and less intermittent as the communicative system matured.

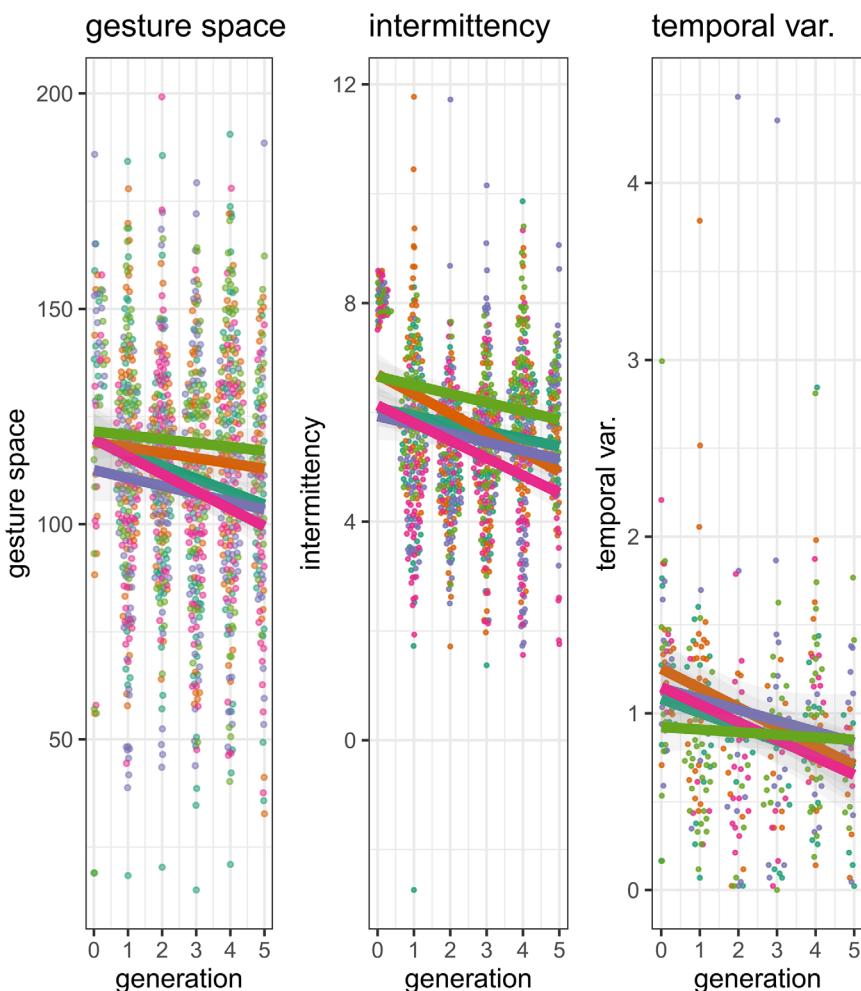


Fig. 7. Change in kinematic properties over generations. Generation trends per chain for intermittency, temporal variability, and gesture space. Over the generations, movements become more smooth (lower intermittency score), show more stable rhythms (lower temporal variability), and more minimized movements (smaller gesture space). There are fewer data points for temporal variability because this can only be computed for comparisons of gestures that have more than two submovements. So temporal variability indicates that *when there was a multi-segmented movement*, then such movements were more rhythmic.

5.2. Network changes over generations

While changes in kinematic complexity suggest an increase in efficiency, they do not by themselves provide evidence of systematicity, another hallmark feature of communicative systems. Here, we assess whether the gesture network as a system shows reduced entropy over generations, which would mean that the interrelationships between gestures become less randomly distributed. Fig. 8a shows that the entropy of gesture networks indeed decreased

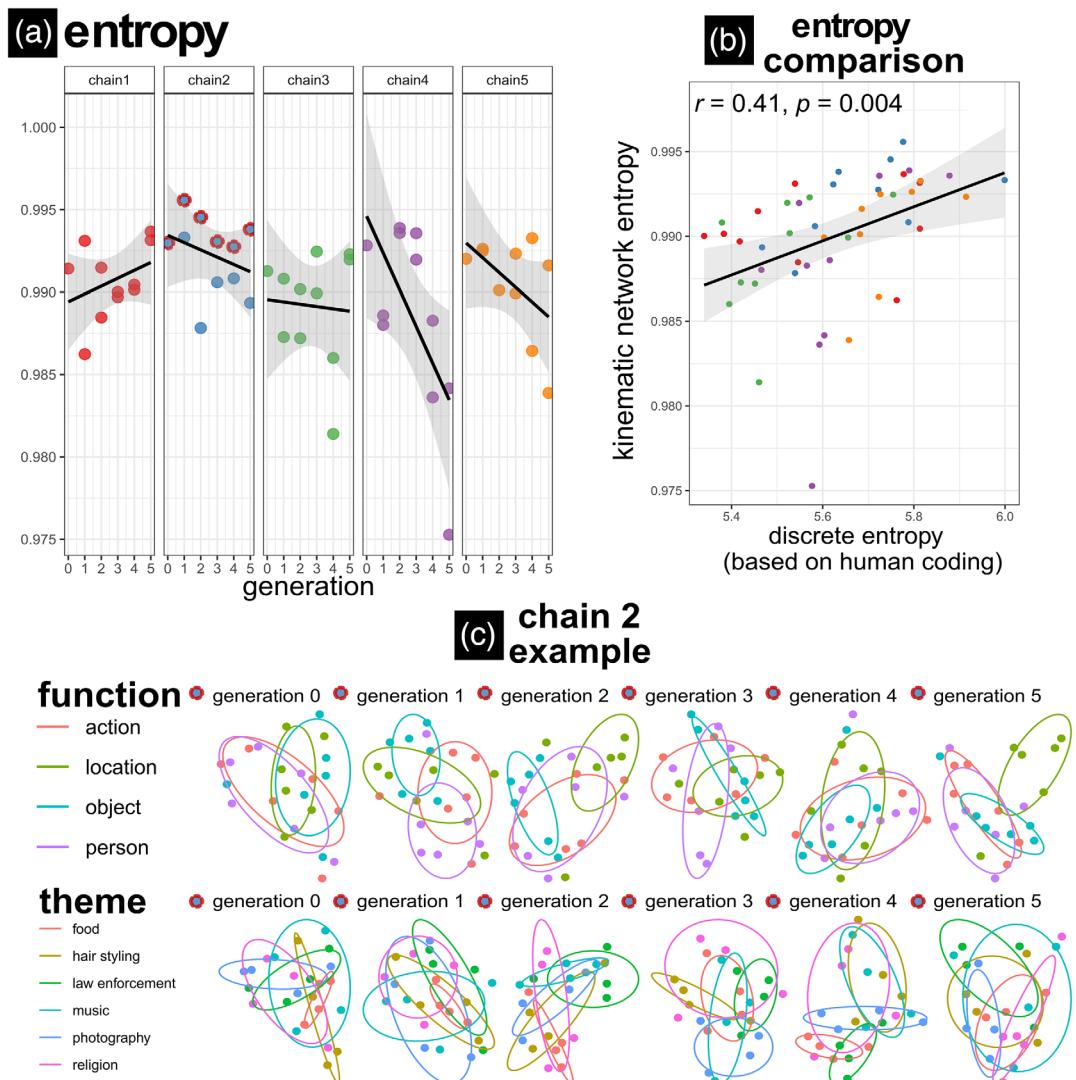


Fig. 8. Changes in networks measures over generations within chains. Panel a shows for each chain the changes over generations in entropy, with generation 0 indicating the seed gesture set. For each generation > 0 , there are two data points as there are two participants in each generation. Entropy tends to decline over the generations, indicating that relationships between tokens became less diverse, possibly indicating systematicity in the way nodes are connected. Panel b shows that the network entropy computed on the kinematic distances was scaled reliably scaling with the discrete entropy computed on human annotated data, which suggests that our kinematic derived measure captures a comparable phenomenon in the evolution of gesture. In panel c, we provide the 2D network representations for each generation in chain 2, color coded and ellipses drawn for theme and function categories for separate rows. We arbitrarily picked one participant for each chain iteration. Note, it is difficult to directly see changes in structure in these representations, which can be directly related to reduction in entropy, but in general, the distances are less diverse over the generations.

as a function of generation in four out of five chains, indicating lower complexity of gesture interrelations as the systems matured. This reduction in entropy, it turns out, scales very reliably with the discrete entropy derived from the manual coding (Fig. 8b).

We tested the trend indicated in Fig. 8a in a mixed-effects regression model similar to the original study (Motamed et al., 2019), with chain as random intercept (random slopes did not converge for these models) and generation as an independent predictor (0–5 generations, with generation 0 being the seed gesture network). Generation was indeed a reliable predictor for network entropy as compared to a base model predicting the overall mean (chi-squared change (1) = 4.75 , $p = .03$, model $R^2 = 0.08$). Model estimates showed that entropy decreased over generations (b estimate = -0.0006 , t (48.00) = -2.19 , $p = .03$, Cohen's $d = -0.63$).

In sum, we find that kinematic network entropy decreases over generations, suggesting a steady increase in systematicity in terms of the distance between gestures in gesture networks. Furthermore, there is a clear scaling relation between the gesture-content entropy computed on human categorical codings and our measure of kinematic entropy, suggesting we are capturing a similar systematic property of these evolving manual languages.

5.3. Relations between kinematic and network properties

So far, we have shown an overall increase in communicative efficiency (as measured by the change in kinematic features over generations) and an increase in systematicity (as measured in decreasing entropy over generations). Reduction in kinematic features may or may not be related to the systematicity found in each gesture chain. Fig. 9 shows how each kinematic property (averaged by the participant) relates to gesture network entropy. We see that network entropy reduces as the average gesture space decreases and movements become less intermittent.

5.4. Chain-unique kinematic evolution over the generations

We have obtained evidence that over generations, kinematic changes go hand-in-hand with a decrease in complexity, and kinematic interrelationships between gestures indicate increasing systematicity. In our final analysis, we use our fine-grained token-level kinematic measures to assess whether unique trajectories in language evolution can be observed. The kinematic changes we have detected could suggest that motoric constraints lead iconic gesture systems to become more similar to each other rather than becoming unique language-like systems. To test whether this is the case, we analyze whether the kinematic interrelationships within a chain diverge from other chains across the study. We predict that even though chains converge in the degree of systematicity (as we found above), they may actually diverge from each other over generations, as conventions are built that are unique to each chain. Fig. 10 shows the main results of this research question.

The network representations at each generation show that early on, gestures do not cluster clearly by chain. However, over generations, especially in generation 3 to 5, we see that gestures from particular chains start to cluster together more prominently, and clusters move away from each other to some extent, indicating greater chain-internal similarity and grow-

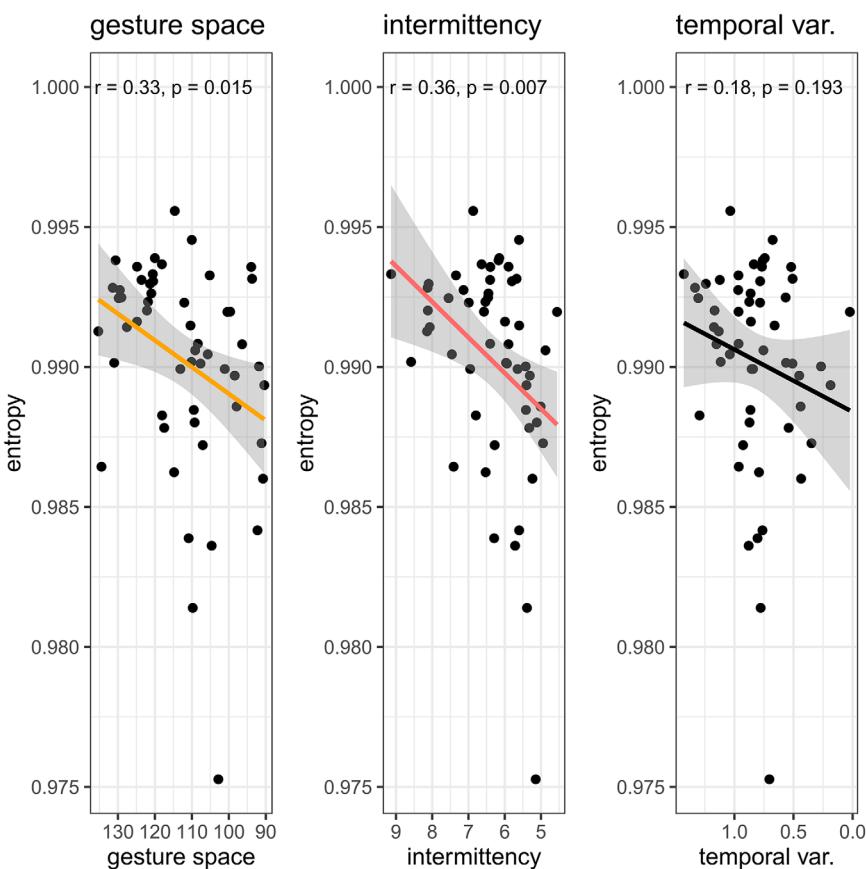


Fig. 9. Relation between kinematic properties and network measures. Correlations for each kinematic property averaged over all tokens for that participant and the concomitant network entropy for that participant. Note that the x-axis is reversed such that reductions in kinematic complexity are related to entropy. The smaller the gesture spaces and the lower the intermittency, the lower the entropy. This indicates that especially reduction in intermittency and gesture space is related with emergence of systematicity on the network level.

ing differences across chains (Fig. 10a). This is indicated by the increase in the Dunn index (Fig. 10b), suggesting that the chain-specific gestures become more compact as a cluster and more removed from other chains. To test this further statistically, we also found that over the generations, Silhouette width increases reliably, $r = .1, p < .001$ (Fig. 10c), suggesting that each chain-specific gesture became more similar to gestures in the same chain and more dissimilar to gestures from the nearest chain in the kinematic space.

In sum, we find that each chain creatively shapes their own gesture system. While our network analysis showed that gesture systems become internally more coherent for all chains, this does not mean that chains become similar to each other. Instead, they forge their own developmental trajectory over generations.

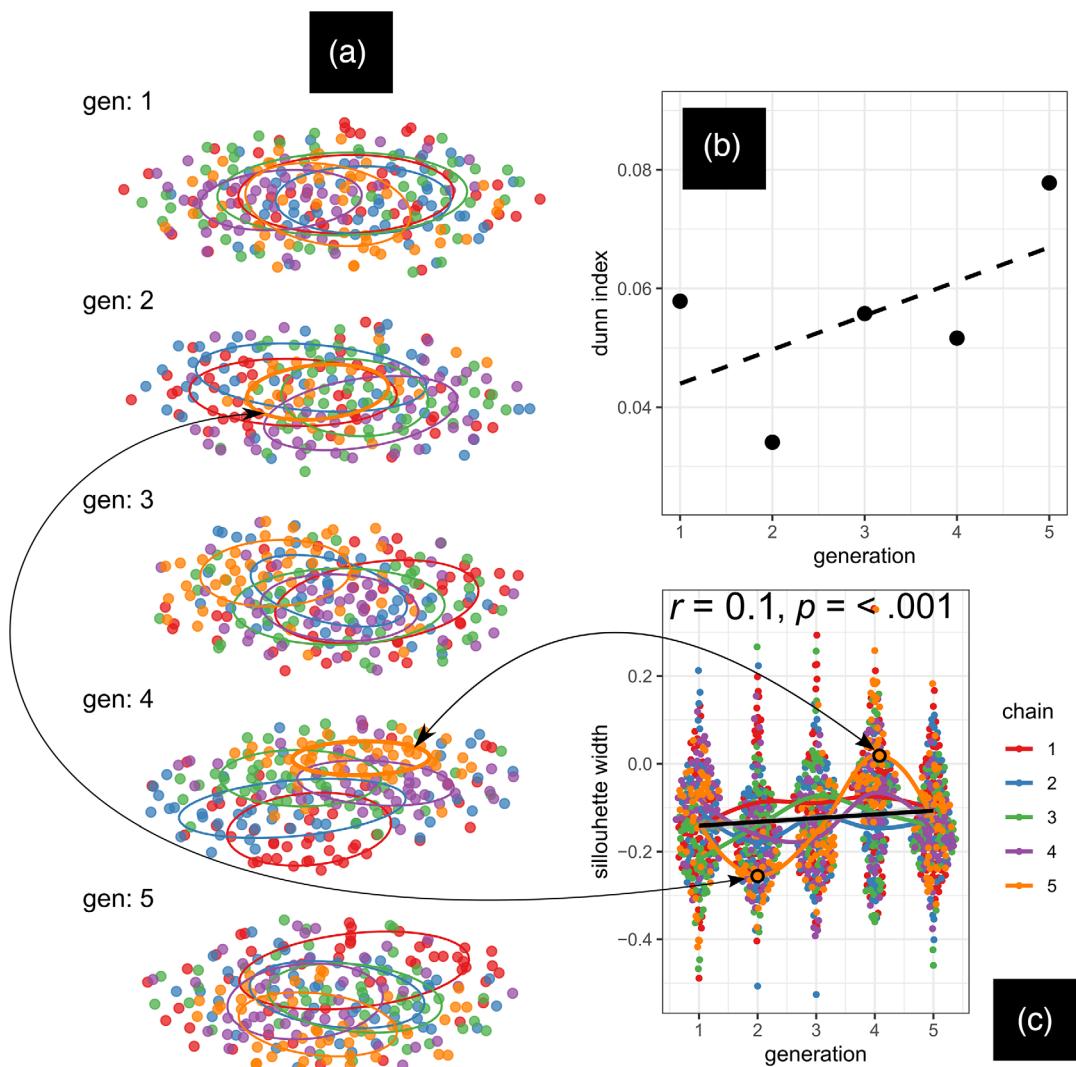


Fig. 10. Evolution of kinematic interrelationships between and within chains over the generations. (a) Network representations of kinematic interrelationships color-coded by chain and with ellipses showing cluster centroids in similarity space. Each point in (a) (and (b)) represents a location in kinematic space or value for a single gesture event produced by a participant that belonged to a certain chain. In generation 1, nodes are scattered and centroids mostly overlap, suggesting gestures do not clearly pattern by chain yet. Over generations, gestures increasingly cluster by chain, and chains drift apart, as seen in the decreasing centroid overlap in generations 3–5. Right panel: For gestures grouped by chains, the Dunn index shows an increase in cluster compactness and distinctiveness over generations (b). Silhouette width reveals at the level of single gestures how the mean distance to tokens from other chains reliably grows relative to the mean distance to tokens from the same chain (c). This is for example apparent when contrasting chain 5 at generation 2 versus that same chain at generation 4, showing higher silhouette scores at generation 4 and more compact clustering in (a).

6. Discussion

Based on signal processing alone, we have detected systematic changes reflective of a linguistically maturing communication system, from continuous multi-articulatory kinematics of silent gestures. We applied computer vision techniques to extract kinematics from video data (e.g., Östling, Börstell, & Courtaux, 2018; Ripperda et al., 2020; Trettenbrein & Zaccarella, 2021). We then quantified kinematic relationships between gestural utterances (Pouw & Dixon, 2019). Our analysis showed that a coding system focusing primarily on the hands would miss important gestural content that can be observed from head movements. We found that including head motion data as opposed to only motion data from the hands allowed us to better distinguish randomly paired gestures from directly related. In line with earlier research suggesting that communicative systems tend toward communicative efficiency (Gibson et al., 2019; Namboodiripad et al., 2016), and going beyond the original study findings, we obtain that gestures reduce their complexity over the generations, by reducing size, submovements, and becoming more rhythmic. Communicative efficiency does not automatically entail systematicity. Therefore, we analyzed gesture networks and showed that communicative tokens have higher systematicity, a finding that is in line with measures derived from manual categorical coding. These results suggest that form analysis can provide information about systematicity, making our method a valuable addition to the toolkit of gesture analysis (where it complements human coding of gesture content) and the cultural evolution of signaling systems more generally, without resorting to categorical semantic coding.

The equivalence of form-based kinematic analysis and content-based manual coding is as much a theoretical advancement as it is a methodological one. It means that aspects of form can directly embody some of the systematic structure inherent in a system (Rączaszek-Leonardi & Kelso, 2008). This is underlined by our final, novel analysis, in which we found that while chains show an increase in systematicity, the gestures that are produced become more distinct with respect to the other chains. Gestures within chains became clearly more clustered over time, meaning their form became more alike within the chain at later generations and more dissimilar from other chains. This is clear evidence of a drift toward chain-specific conventions, suggesting the emergence of gesture kinematic dialects.

In sum, the current analysis provides new insights into how gesture movements linguistically evolve. First, increased communicative efficiency was previously judged absent on the basis of gestural unit length, but we show it is present in the kinematics on relevant dimensions as salience, segmentation, and temporality. Second, kinematic properties change coherently to simplify in structure, which it turns out, is relatable to the simplification of the semantic content of gestures, a non-trivial equivalence given the fact that form and content are generally forcefully distinguished. Finally, we obtain that kinematic dialects emerge, which goes beyond the goals of the original study altogether.

6.1. Shortcomings and advantages

There are two caveats to the analyses presented here. First, kinematic analysis alone cannot say anything precise enough to determine the semiotic content of tokens (cf. Pouw et al.,

2021). Therefore, understanding the semiotic content of human communication will always require extensive human analysis (Sandler, 2018). Still, the kinematic analysis provides us with a unique grasp of aspects of rich evolving communicative systems that may elude human coders or would be too resource-demanding to manually code. The most productive way forward, therefore, is to use methodological triangulation by combining kinematic measures and semiotic analyses. For instance, kinematic measures would allow the detection of subtle changes in gesture form or system-level structure over time. This analysis can then be subjected to linguistic analysis or human coding to understand semiotic and structural aspects of the evolving system. For example, we find that the inclusion of head movements meaningfully improves our kinematic analysis. While the original experiment focused on the hands as the most important articulators (Motamed et al., 2019), our finding invites consideration of how head movements may be recruited as part of a culturally evolving semiotic system, a finding that has implications beyond purely methodological concerns. The general picture that emerges is that kinematic measures allow a bottom-up, data-driven approach to be informed and enriched by qualitative analysis solicited by first quantitatively identifying the “active regions” of the data.

A second caveat concerns a limitation to DTW. One finding from prior work on the evolution of compositionality is that holistic gestures may become segmented, freeing up individual elements to be recombined, possible in different orders. Now, consider two gestures that contain identical segments in different orders. While human coders would likely recognize the commonality and judge these as highly similar, the DTW algorithm is sensitive to the ordering and would judge them as very different. Indeed, in the animated gesture network on our supplemental page (<https://osf.io/wbmf9/>), it can be seen that though “handcuffs” and “prison” share a gesture component, they are deemed quite dissimilar according to our procedure due to different ordering. So our DTW analysis may at times judge sequences of gestures highly dissimilar when in fact they are merely ordered differently. There are ways to circumvent this by only looking for trajectory overlaps rather than ordering through time (Pouw & Dixon, 2019), but such analyses go beyond the current approach. Importantly, this limitation stacks the deck against finding increases in systematicity, so it speaks to the robustness of our measures (and perhaps indicates the limited occurrence of reordering elements) that some degree of systematicity is nonetheless recovered.

Both of these caveats mean that our approach to kinematics, like all quantitative analyses of human behavior, requires some degree of human oversight (for meaningful implementation) and human insight (for judicious interpretation). When these requirements are met, we believe that our fully reproducible and automatable methods can make important contributions to the systematic study of continuous communicative signals.

6.2. *Embodied language evolution*

In our view, the current findings underline an understanding of language evolution as an embodied process. Consider, for example, our finding that gesture tokens become simpler in several dimensions: smaller size, fewer submovements, and less temporal variability. This simplification seems to be a reduction in articulatory effort. Making a smoother, smaller, and

more rhythmic movements reduces the states that a sensorimotor routine needs to visit (Kelso, Tuller, Vatikiotis-Bateson, & Fowler, 1984). Thus, communicative efficiency increased over the generations at the motoric level, an effect not captured by the content-level repetition and gesture length measure used in the original study (Motamedi et al., 2019). Note that kinematic efficiency could also potentially increase the learnability and comprehensibility of the gestures. Speech perception in noisy conditions is more optimal when speech is more rhythmic (Wang, Kong, Zhang, Wu, & Li, 2018). We submit that the optimization of sensorimotor routines of communication is an integral part of the increased efficiency of communication. This is fully in line with early insights on the psychobiology of language by Zipf or MacNeilage (MacNeilage, 2010; Zipf, 1935), which have been overshadowed to some extent by the focus on codes and information content in the wake of Shannon. Our measures show how it is possible to unite fine-grained measures of the biomechanics of evolving continuous signals with information-theoretic notions like system-level entropy. Low-level action-perception optimization and code-level systematic optimization should not be treated as categorically distinct processes.

We find a similar reduction of complexity of pronunciation as in novice learners of American Sign Language (ASL). ASL learners spatially reduce their signs as they become more fluent (Lupton & Zelaznik, 1990; Wilbur, 1990). Compound signs also become shorter as multi-component signs merge into efficiently produced single signs, in a way that is mirrored in the data we have analyzed here. As sign systems evolve, suboptimal organizations of sub-movements give way to more efficient signs and temporally extended sequences become more coordinated (Bernstein, 1967; Kelso, Tuller, & Harris, 1983). This makes the dynamics of the evolving gesture systems we have studied here similar in some ways to those of full-blown manual language systems, such as ASL.

Another sign of the maturity of the culturally evolved gestural systems we have studied here is the fact that head movements can function as cues reliably distinguishing paired gestures. This finding resonates with the known grammatic, phonetic, and prosodic functions that head movements have in sign languages such as ASL (Tyrone & Mauk, 2016). Indeed, as Sandler (2018) has argued for sign languages, the expressive power of the body lies in the combination of different articulators, combined into a single synergetic utterance.

6.3. Future directions

While we have focused on silent gesture as a test case here, our analyses are applicable to any continuous signaling system. Similar approaches are applied in animal signaling systems, where high dimensional features of tokens are mapped onto lower-dimensional space to identify stably distinct patterns (Sainburg, Thielk, & Gentner, 2019). Staying within the domain of human communication, the current analysis could be applied to speech acoustics, semantics, and gesture kinematics in unison (Pouw & Dixon, 2019; Pouw et al., 2021). Indeed, in a study by Perlman, Dale, and Lupyán (2015), it is shown how dynamic aspects of vocalization signaling systems become more efficient, similar to our current reduction in kinematic complexity. These findings, together with work showing the tight connection between speech and gesture (Bosker & Peeters, 2020; Pouw, de Jonge-Hoekstra, Harrison, Paxton, & Dixon, 2020; Pouw,

Harrison, Esteve-Gibert, & Dixon, 2020), make it a natural next step to look at multimodal iterated learning experiments. Furthermore, our approach can inform work on communicative alignment in conversations (Rasenberg, Özyürek, & Dingemanse, 2020) or the ways in which people can repeat aspects of each other's communicative behavior. In short, our analysis here has not only yielded new findings in the cultural evolution of gestural communication systems but opens the door toward a broader research program in which the action-perception aspects of communication systems are studied alongside their structural and semiotic aspects.

7. Conclusion

Human communicative behavior tends to combine categorical elements and continuous properties, but for technological as well as theoretical reasons, the categorical elements of evolving linguistic systems have long received more attention than their continuous aspects. Here, we have contributed to the study of multimodal language and cognition by considering the gesture kinematics of evolving gestural systems. We have used computer vision techniques to analyze the kinematic properties of evolving gestural systems, showing that over generations of learners, the dynamics of head and upper limb movements become simpler, increase in systematicity, and give rise to kinematic dialects. Our kinematic measures help characterize fine-grained levels of linguistic organization that remain out of reach of content-based discretized coding approaches, providing novel insights that corroborate and complement prior approaches. Our findings provide an unprecedented view of how gestures become structured and increasingly language-like as they evolve, in ways that are directly related to the coordination and simplification of bodily movements. While considerations of communicative efficiency and systematicity have so far been mostly based on analyses of discrete symbol systems like written words and text corpora, our work shows how hallmark features of linguistic systems may be grounded directly in the biomechanical properties of dynamically evolving systems of continuous signals.

Acknowledgments

This work is supported by a Donders Fellowship awarded to Wim Pouw (supervised by Ashlı Özyürek) and is supported by the Language in Interaction consortium project “Communicative Alignment in Brain & Behavior” (CABB). This manuscript has been written in Rmarkdown, a code-embedded version can be found on our OSF page <https://osf.io/56jwb>.

Open Research Badges



This article has earned Open Data and Open Materials badges. Data and materials are available at <https://osf.io/56jwb> and <https://datashare.ed.ac.uk/handle/10283/3191>.

References

- Beecks, C., Hassani, M., Brenger, B., Hinnell, J., Schüller, D., Mittelberg, I., & Seidl, T. (2016). Efficient query processing in 3D motion capture gesture databases. *International Journal of Semantic Computing*, 10(01), 5–25. <https://doi.org/10.1142/S1793351X16400018>
- Beecks, C., Hassani, M., Hinnell, J., Schüller, D., Brenger, B., Mittelberg, I., & Seidl, T. (2015). Spatiotemporal similarity search in 3D motion capture gesture streams. In C. Claramunt, M. Schneider, R. C. -W. Wong, L. Xiong, W. -K. Loh, C. Shahabi & K. -J. Li (Eds.), *Advances in spatial and temporal databases* (pp. 355–372). Cham: Springer International Publishing. https://doi.org/10.1007/978-3-319-22363-6_19
- Bernstein, N. (1967). *The co-ordination and regulations of movements* (1st English ed.). Oxford, England: Pergamon Press.
- Bolinger, D. (1968). *Aspects of language*. New York: Harcourt, Brace, and World.
- Börstell, C., & Lepic, R. (2020). Spatial metaphors in antonym pairs across sign languages. *Sign Language & Linguistics*, 23(1-2), 112–141. <https://doi.org/10.1075/sll.00046.bor>
- Bosker, H. R., & Peeters, D. (2020). Beat gestures influence which speech sounds you hear. *Proceedings of the Royal Society B: Biological Sciences*, 288(1946), 20202419. <http://doi.org/10.1098/rspb.2020.2419>
- Cao, Z., Simon, T., Wei, S. -E., & Sheikh, Y. (2017). Realtime multi-person 2D pose estimation using part affinity fields. *2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)* Honolulu, HI (pp. 1302–1310). <https://doi.org/10.1109/CVPR.2017.143>
- Christiansen, M. H., & Chater, N. (2016). The now-or-never bottleneck: A fundamental constraint on language. *Behavioral and Brain Sciences*, 39, e62. <https://doi.org/10.1017/S01402525X1500031X>
- Cornish, H., Dale, R., Kirby, S., & Christiansen, M. H. (2017). Sequence memory constraints give rise to language-like structure through iterated learning. *PLoS ONE*, 12(1), e0168532. <https://doi.org/10.1371/journal.pone.0168532>
- Csárdi, G. (2019). Package ‘igraph’ network analysis and visualization (Version 1.2.4.1). Retrieved from <http://bioconductor.statistik.tu-dortmund.de/cran/web/packages/igraph/igraph.pdf>
- Ćwiek, A., Fuchs, S., Draxler, C., Asu, E. L., Dedić, D., Hiovain, K., ... & Perlman, M. (2021). Novel vocalizations are understood across cultures. *Scientific Reports*, 11(1), 1–12.
- Dale, R., & Kello, C. T. (2018). “How do humans make sense?” Multiscale dynamics and emergent meaning. *New Ideas in Psychology*, 50, 61–72. <https://doi.org/10.1016/j.newideapsych.2017.09.002>
- Dingemanse, M., Blasi, D. E., Lupyan, G., Christiansen, M. H., & Monaghan, P. (2015). Arbitrariness, iconicity, and systematicity in language. *Trends in Cognitive Sciences*, 19(10), 603–615. <https://doi.org/10.1016/j.tics.2015.07.013>
- Donaldson, J. (2016). tsne: T-distributed stochastic neighbor embedding for R (t-SNE) (Version 0.1-3). Retrieved from <https://CRAN.R-project.org/package=tsne>
- Eagle, N., Macy, M., & Claxton, R. (2010). Network diversity and economic development. *Science*, 328(5981), 1029–1031. <https://doi.org/10.1126/science.1186605>
- Enfield, N. J. (2009). *The anatomy of meaning: Speech, gesture, and composite utterances*. Cambridge: Cambridge University Press.
- Enfield, N. J. (2014). *Natural causes of language: Frames, biases, and cultural transmission*. Berlin: Language Science Press.
- Ergin, R., Kürşat, L., Hartzell, E., & Jackendoff, R. (2021, January 18). Central taurus sign language: On the edge of conventionalization. <https://doi.org/10.31234/osf.io/x9emd>
- Gerwing, J., & Bavelas, J. (2004). Linguistic influences on gesture’s form. *Gesture*, 4(2), 157–195. <https://doi.org/10.1075/gest.4.2.04ger>
- Gibson, E., Futrell, R., Piantadosi, S. P., Dautriche, I., Mahowald, K., Bergen, L., & Levy, R. (2019). How efficiency shapes human language. *Trends in Cognitive Sciences*, 23(5), 389–407. <https://doi.org/10.1016/j.tics.2019.02.003>
- Giorgino, T. (2009). Computing and visualizing dynamic time warping alignments in R: The dtw package. *Journal of Statistical Software*, 1(7), 1–25. <https://doi.org/10.18637/jss.v031.i07>

- Haviland, J. B. (2013). The emerging grammar of nouns in a first generation sign language: Specification, iconicity, and syntax. *Gesture*, 13(3), 309–353. <https://doi.org/10.1075/gest.13.3.04hav>
- Hogan, N., & Sternad, D. (2009). Sensitivity of smoothness measures to movement duration, amplitude and arrests. *Journal of Motor Behavior*, 41(6), 529–534. <https://doi.org/10.3200/35-09-004-RC>
- Holler, J., & Wilkin, K. (2011). An experimental investigation of how addressee feedback affects co-speech gestures accompanying speakers' responses. *Journal of Pragmatics*, 43(14), 3522–3536. <https://doi.org/10.1016/j.pragma.2011.08.002>
- Kelso, J. A. S., Tuller, B., & Harris, K. (1983). A "Dynamic Pattern" perspective on the control and coordination of movement. In P. F. MacNeilage (Ed.), *The production of speech* (pp. 137–173). Berlin: Springer-Verlag. Retrieved from https://link.springer.com/chapter/10.1007/978-1-4613-8202-7_7
- Kelso, J. A. S., Tuller, B., Vatikiotis-Bateson, E., & Fowler, C. A. (1984). Functionally specific articulatory cooperation following jaw perturbations during speech: Evidence for coordinative structures. *Journal of Experimental Psychology. Human Perception and Performance*, 10(6), 812–832.
- Kendon, A. (2004). *Gesture: Visible action as utterance*. Cambridge: Cambridge University Press.
- Kirby, S., Griffiths, T., & Smith, K. (2014). Iterated learning and the evolution of language. *Current Opinion in Neurobiology*, 28, 108–114. <https://doi.org/10.1016/j.conb.2014.07.014>
- Lepic, R., Börstell, C., Belsitzman, G., & Sandler, W. (2016). Taking meaning in hand: Iconic motivations in two-handed signs. *Sign Language & Linguistics*, 19(1), 37–81. <https://doi.org/10.1075/sll.19.1.02lep>
- Lupton, L. K., & Zelaznik, H. N. (1990). Motor learning in sign language students. *Sign Language Studies*, 1067(1), 153–174. <https://doi.org/10.1353/sls.1990.0020>
- MacNeilage, P. F. (2010). *The origin of speech*. Oxford, England: Oxford University Press.
- McNeill, D. (2005). *Gesture and thought*. Chicago: The University of Chicago Press.
- Motamedi, Y., Schouwstra, M., Smith, K., Culbertson, J., & Kirby, S. (2019). Evolving artificial sign languages in the lab: From improvised gesture to systematic sign. *Cognition*, 192, 103964. <https://doi.org/10.1016/j.cognition.2019.05.001>
- Mueen, A. K., & Keogh, E. (2016). Extracting optimal performance from dynamic time warping. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, San Francisco, CA (pp. 2129–2130). <https://doi.org/10.1145/2939672.2945383>
- Muller, M. (2007). *Information retrieval for music and motion*. Heidelberg, Germany: Springer.
- Namboodiripad, S., Lenzen, D., Lepic, R., & Verhoef, T. (2016). Measuring conventionalization in the manual modality. *Journal of Language Evolution*, 1(2), 109–118. <https://doi.org/10.1093/jole/lzw005>
- Ortega, G., & Özyürek, A. (2019). Types of iconicity and combinatorial strategies distinguish semantic categories in silent gesture. *Language and Cognition*, 12(1), 84–113. <https://doi.org/10.1017/langcog.2019.28>
- Östling, R., Börstell, C., & Courtaux, S. (2018). Visual iconicity across sign languages: Large-scale automated video analysis of iconic articulators and locations. *Frontiers in Psychology*, 9, 725. <https://doi.org/10.3389/fpsyg.2018.00725>
- Padden, C. A., Meir, I., Hwang, S. -O., Lepic, R., Seegers, S., & Sampson, T. (2013). Patterned iconicity in sign language lexicons. *Gesture*, 13(3), 287–308. <https://doi.org/10.1075/gest.13.3.03pad>
- Pattee, H. H., & Račzaszek-Leonardi, J. (2012). *LAWS, LANGUAGE and LIFE: Howard Pattee's classic papers on the physics of symbols with contemporary commentary*. Dordrecht: Springer Science & Business Media. Retrieved from <http://books.google.com?id=raEQodcVdYQC>
- Perlman, M., Dale, R., & Lupyán, G. (2015). Iconicity can ground the creation of vocal symbols. *Royal Society Open Science*, 2(8), 150152. <https://doi.org/10.1098/rsos.150152>
- Pouw, W., de Jonge-Hoekstra, L., Harrison, S. J., Paxton, A., & Dixon, J. A. (2020). Gesture-speech physics in fluent speech and rhythmic upper limb movements. *Annals of the New York Academy of Sciences*, 1491(1), 89–105. <https://doi.org/10.1111/nyas.14532>
- Pouw, W., & Dixon, J. A. (2019). Gesture networks: Introducing dynamic time warping and network analysis for the kinematic study of gesture ensembles. *Discourse Processes*, 57(4), 301–319. <https://doi.org/10.1080/0163853X.2019.1678967>

- Pouw, W., Harrison, S. J., Esteve-Gibert, N., & Dixon, J. A. (2020). Energy flows in gesture-speech physics: The respiratory-vocal system and its coupling with hand gestures. *The Journal of the Acoustical Society of America*, 148(3), 1231–1247. <https://doi.org/10.1121/10.0001730>
- Pouw, W., & Trujillo, J. P. (2019). Materials Tutorial Gespin2019—Using video-based motion tracking to quantify speech-gesture synchrony. Retrieved from 10.17605/OSF.IO/RXB8J
- Pouw, W., Wit, J., Bögels, S., Rasenberg, M., Milivojevic, B., & Özyürek, A. (2021). Semantically related gestures move alike: Towards a distributional semantics of gesture kinematics. *Proceedings of the 23rd International Conference on Human-Computer Interaction*, Washington, DC. <https://doi.org/10.31219/osf.io/pgq6m>
- Rączaszek-Leonardi, J., & Kelso, S. J. A. (2008). Reconciling symbolic and dynamic aspects of language: Toward a dynamic psycholinguistics. *New Ideas in Psychology*, 26(2), 193–207. <https://doi.org/10.1016/j.newideapsych.2007.07.003>
- Rasenberg, M., Özyürek, A., & Dingemanse, M. (2020). Alignment in multimodal interaction: An integrative framework. *Cognitive Science*, 44(11), e12911. <https://doi.org/10.1111/cogs.12911>
- Ravignani, A., Delgado, T., & Kirby, S. (2016). Musical evolution in the lab exhibits rhythmic universals. *Nature Human Behaviour*, 1, 0007. <https://doi.org/10.1038/s41562-016-0007>
- Ripperda, J., Drijvers, L., & Holler, J. (2020). Speeding up the detection of non-iconic and iconic gestures (SPUD-NIG): A toolkit for the automatic detection of hand movements and gestures in video data. *Behavior Research Methods*, 52(4), 1783–1794. <https://doi.org/10.3758/s13428-020-01350-2>
- Sainburg, T., Thielk, M., & Gentner, T. Q. (2019). Latent space visualization, characterization, and generation of diverse vocal communication signals. *bioRxiv*, 870311, [https://doi.org/10.1101/870311 bioRxiv](https://doi.org/10.1101/870311)
- Sandler, W. (2018). The body as evidence for the nature of language. *Frontiers in Psychology*, 9, 1782. <https://doi.org/10.3389/fpsyg.2018.01782>
- Sato, A., Schouwstra, M., Flaherty, M., & Kirby, S. (2020). Do all aspects of learning benefit from iconicity? Evidence from motion capture. *Language and Cognition*, 12(1), 36–55. <https://doi.org/10.1017/langcog.2019.37>
- Scott-Phillips, T. C., & Kirby, S. (2010). Language evolution in the laboratory. *Trends in Cognitive Sciences*, 14(9), 411–417. [10.1016/j.tics.2010.06.006](https://doi.org/10.1016/j.tics.2010.06.006)
- Senghas, A., Kita, S., & Özyürek, A. (2004). Children creating core properties of language: Evidence from an emerging sign language in nicaragua. *Science*, 305(5691), 1779–1782. <https://doi.org/10.1126/science.1100199>
- Trettenbrein, P. C., & Zaccarella, E. (2021). Controlling video stimuli in sign language and gesture research: The OpenPoseR package for analyzing OpenPose motion-tracking data in R. *Frontiers in Psychology*, 12, 628728. <https://doi.org/10.3389/fpsyg.2021.628728>
- Trujillo, J. P., Simanova, I., Bekkering, H., & Özyürek, A. (2019). The communicative advantage: How kinematic signaling supports semantic comprehension. *Psychological Research*, 84(7), 1897–1911.
- Trujillo, J. P., Vaitonyte, J., Simanova, I., & Özyürek, A. (2019). Toward the markerless and automatic analysis of kinematic features: A toolkit for gesture and movement research. *Behavior Research Methods*, 51(2), 769–777. <https://doi.org/10.3758/s13428-018-1086-8>
- Tyrone, M. E., & Mauk, C. E. (2016). The phonetics of head and body movement in the realization of American sign language signs. *Phonetica*, 73(2), 120–140. <https://doi.org/10.1159/000443836>
- van der Maaten, L., & Hinton, G. (2008). Visualizing data using t-SNE. *Journal of Machine Learning Research*, 9(86), 2579–2605. Retrieved from <http://jmlr.org/papers/v9/vandermaaten08a.html>
- Verhoef, T., Kirby, S., & de Boer, B. (2014). Emergence of combinatorial structure and economy through iterated learning with continuous acoustic signals. *Journal of Phonetics*, 43, 57–68. <https://doi.org/10.1016/j.wocn.2014.02.005>
- Verhoef, T., Kirby, S., & de Boer, B. (2016). Iconicity and the emergence of combinatorial structure in language. *Cognitive Science*, 40(8), 1969–1994. <https://doi.org/10.1111/cogs.12326>
- Wang, M., Kong, L., Zhang, C., Wu, X., & Li, L. (2018). Speaking rhythmically improves speech recognition under “cocktail-party” conditions. *The Journal of the Acoustical Society of America*, 143(4), EL255. <https://doi.org/10.1121/1.5030518>

- Wilbur, R. B. (1990). An experimental investigation of stressed sign production. *International Journal of Sign Linguistics*, 1(1), 41–60.
- Yadav, A. K., Tomar, D., & Agarwal, S. (2013). Clustering of lung cancer data using Foggy K-means. *2013 International Conference on Recent Trends in Information Technology (ICRTIT)*, Chennai, India (pp. 13–18). <https://doi.org/10.1109/ICRTIT.2013.6844173>
- Zipf, G. K. (1935). *The psycho-biology of language*. Boston: Houghton Mifflin.

Supporting Information

Additional supporting information may be found online in the Supporting Information section at the end of the article.

Supplementary Material